

Introduction to OS

Lecture 1

Overview

■ Operating Systems basic concepts:

- What is OS?
- Brief History
 - Motivation for OS
- Overview of Modern OSes

■ Operating System Structures

- OS components
- Types of kernel

■ Virtual Machines

What is OS?


- Incorrect/Incomplete definition:

- It is the desktop when you boot up your PC
 - The "thing" that stores your games
 - Windows! (or Mac!) (or Linux!)

- One simple definition:

- A **program** that acts as an **intermediary** between a **computer user** and the **computer hardware**
 - **Wikipedia:** An operating system (OS) is system software that manages computer hardware, software resources, and provides common services for computer programs.

Illustration: What is an OS?



- A simplified view:
 - Will be refined as we move along
- The most general version:
 - Hardware (not only computer!)
 - User (can be application programs or actual person!)

Example of Common OS

- On Computer:
 - Windows 11 / 10, ...
 - MacOS
 - Linux distros: Ubuntu, Redhat, Debian,
- On Smartphone:
 - iOS, Android
- Other hardware with OS:
 - Game console: PS5, Xbox, Nintendo Switch, Steam Machine
 - Home appliance: Smart TV, Smart Watch, Smart Appliances ...

To invent the future, you must understand the past


BRIEF HISTORY OF OS

Brief History of OS

- Essentially, OS evolves with:
 - Computer hardware
 - User application and usage pattern
- The "first" computers:
 - Electronic Numerical Integrator And Computer (**ENIAC**)
 - 1945
 - Program controlled by cables and switches
 - Harvard Mark I:
 - 1944
 - Program controlled by punched paper tape

OS for the first computers

- OS Type:
 - NO OS




- Programs directly interact with hardware
 - Reprogram by changing **physical** configuration of hardware
- **Advantage:**
 - Minimal overhead
- **Disadvantage:**
 - Not portable
 - Inefficient use of computer!

Mainframes: The "Big Iron"

- Commonly used by large corporations in 60s, 70s
- Common features:
 - No interactive interface
 - Accept programs in the form of:
 - Paper tape, magnetic tape, punch card
 - Support batch processing only
 - Very costly
 - Usually "rented" instead of owned
- Example:
 - IBM 360
 - Cost 5 billion US dollar in 1964

OS for Mainframes

- OS Type:
 - **Batch OS**



- Batch OS:
 - Execute user program (a.k.a ***job***) one at a time
 - Load job from media, execute, collect result
- User Job:
 - Still interact with hardware directly
 - With additional information for the OS
 - Resource required
 - Job specification


OS for Mainframes: Improvements

- Simple batch processing is inefficient:
 - CPU idle when perform I/O
- One possible Improvements:
 - Multiprogramming:
 - loads multiple jobs and runs other jobs when I/O needs to be done
 - Overlaps computation with I/O
- Another development of OS during this period (70s):
 - Time-Sharing OS

Time-Sharing OS

- Features:
 - Allow multiple users to interact with machine using terminals (teletypes)
 - User job scheduling
 - Illusion of Concurrency
 - Memory management
- Famous Examples:
 - CTSS developed at MIT 1960s
 - Multics (1970s)
 - Considered as the parent of Unix
 - Pushed the state of art in virtual memory, security
- Not so different from using Unix servers today but more primitive

Time-sharing OS: Illustration



- OS manages the sharing of:
 - CPU time, memory and storage
- ***Virtualization*** of hardware:
 - Each program executes **as if** it has all the resources to itself

Minicomputer and Unix

- Minicomputer follows the mainframe:
 - A "mini" version of mainframe:
 - Smaller and cheaper
 - Example:
 - Digital Equipment Corp (DEC) PDP-11
- Famous OS:
 - Unix
 - Developed by AT&T employees, including Ken Thompson, Dennis Ritchie, Douglas McIlroy, and Joe Ossanna
 - Ken Thompson and Dennis Ritchie
 - Invented the C programming language as well!!

Personal Computer

- Apple II PC:

- First successfully produced mass home computer
 - Designed by Steve Wozniak (alone!)

- IBM PC:

- The first generic PC
 - PC becoming truly a collection of commodity hardware components
 - Leads to dominance of Microsoft OS in PCs: MSDOS then Windows

OS on Personal Computer

- Machine (can be) dedicated to user, not timeshared between multiple users
 - Give rise to **personal OS**
- Several Models:
 - Windows model:
 - Single user at a time but possibly more than 1 user can access
 - Dedicated machine
 - Unix model:
 - One user at the workstation but other users can access remotely
 - General time sharing model

Why do we need OS?

MOTIVATIONS OF OS

Motivation for OS: Abstraction

- Large variation in hardware configurations
- Example (Hard disk):
 - Different capacity (500MB, 320GB, 1.5TB etc)
 - Different capabilities:
 - Rotation per minutes (RPM)
 - Access (read/write) speed
 - Etc
- However, hardware in the same category has well defined and common functionality
 - Example (Hard disk): Store and retrieve information

Motivation for OS: Abstraction

- Operating System serves as an **abstraction**:
 - Hide the different low level details
 - Present the common high level functionality to user
- The user can then perform essential tasks **through** operating system
 - No need to concern with low level details
- Provides:
 - Efficiency, programmability and portability

Motivation for OS: Resource Allocator

- Program execution requires multiple resources:
 - CPU, memory, I/O devices etc
- For better utilization of resources, multiple programs should be allowed to execute simultaneously
- OS is a **resource allocator**
 - Manages all resources
 - CPU, Memory, Input / Output devices
 - Arbitrate potentially conflicting requests
 - for efficient and fair resource use

Motivation for OS: Control Program

- Program can misuse the computer:
 - Accidentally: due to coding bugs
 - Maliciously: virus, malware etc
- Multiple users can share the computer:
 - Tricky to ensure separate user space
- OS is a **control program**
 - Controls execution of programs
 - prevent errors and improper use of the computer
 - Provides security and protection


Motivation for OS: Summary

- Manage resources and coordination
 - process synchronization, resource sharing
- Simplify programming
 - abstraction of hardware, convenient services
- Enforce usage policies
- Security and protection
- User Program Portability:
 - Across different hardware
- Efficiency
 - Sophisticated implementations
 - Optimized for particular usage and hardware

The families of modern OS

OVERVIEW OF MODERN OS

Modern OS: Overview




Common Architecture for OS

OS STRUCTURE

Operating System Structures


- We have identified the major capabilities of an OS
 - i.e. the *specification* of the OS
- Let us now consider:
 - The best way to provides these capabilities
 - i.e. the *implementations* of the OS
- Operating system structure:
 - *Organization* of the various components
 - Important factors:
 - Flexibility
 - Robustness
 - Maintainability

Illustration: High level view of OS



- Operating System is essentially a software
 - Runs in ***kernel mode***: Have complete access to all hardware resources
- Other software executes in ***user mode***
 - With limited (or controlled) access to hardware resources

Illustration: Generic OS Components



- **A:** OS executing machine instructions
- **B:** normal machine instructions executed (program/library code)
- **C:** calling OS using ***system call interface***
- **D:** user program calls library code
- **E:** system processes
 - Provide high level services, usually part of OS

OS as a Program

- OS is also known as the **kernel**
 - Just another program with some special features
 - Deals with hardware issues
 - Provides system call interface
 - Special code for interrupt handlers, device drivers
- Kernel code has to be different than normal programs:
 - Cannot use system calls in kernel code
 - Cannot use normal libraries
 - No “normal” I/O
- Consider this:
 - Normal programs use OS: what does OS use? ☺

Implementing Operating System

■ Programming Language:

- Historically in assembly/machine code
- Now in *HLLs*:
 - Especially C/C++
- Heavily hardware architecture dependent

■ Common code organization:

1. Machine independent HLL
2. Machine dependent HLL
3. Machine dependent assembly code

■ Challenges:

- “No one else” to rely on for nice services
- Debugging is hard
- Complexity
- Enourmous Codebase


OS Structures

- Several ways to structure an OS:
 - **Monolithic**
 - **Microkernel**
 - Layered
 - Client-Server
 - Exokernel
 - etc
- We will cover the first two in details:
 - They represent the whole range of possibilities
 - Most other approaches are variant or improvement

Monolithic OS

- Kernel is:
 - One **BIG** special program
 - Various services and components are integral part
 - Good SE principles is still possible with:
 - modularization
 - separation of interfaces and implementation
- This is the traditional approach taken by:
 - Most Unix variants, Windows NT/XP
- **Advantages:**
 - Well understood
 - Good performance
- **Disadvantages:**
 - Highly coupled components
 - Usually devolved into very complicated internal structure

Monolithic Kernel Illustration




Generic Architecture of Monolithic OS Components

Microkernel OS

- Kernel is:
 - Very small and clean
 - Only provides basic and essential facilities:
 - Inter-Process Communication (IPC)
 - Address space management
 - Thread management
 - etc
- Higher level services:
 - Built on top of the basic facilities
 - Run as server process outside of the OS
 - Use IPC to communicate
- **Advantages:**
 - Kernel is generally more robust and more extendible
 - Better isolation and protection between kernel and high level services
- **Disadvantages:**
 - Lower Performance

Microkernel Components



Other Operating System Structure

■ Layered Systems:

- Generalization of monolithic system
- Organize the components into hierarchy of layers
 - Upper layers make use of the lower layers
 - Lowest layer is the hardware
 - Highest layer is the user interface

■ Client-Server Model

- Variation of microkernel
- Two classes of processes:
 - Client process request service from server process
 - Server Process built on top of the microkernel
 - Client and Server process can be on separate machine!

Ways of running OSes

VIRTUAL MACHINES


Motivation: Why Virtual Machines

- OS assumes total control of the hardware:
 - What if we want to run several OSes on the same hardware at the same time?
 - Example: Cloud Computing (**IaaS**, Infrastructure as a Service)
- OS is hard to debug / monitor:
 - How do we observe the working of the OS?
 - How do we test a potentially destructive implementation?

Definition: Virtual Machine


- Virtual Machine:
 - A software emulation of hardware
 - **Virtualization** of underlying hardware
 - Illusion of complete hardware to level above: memory, CPU, hard disk etc...
 - Normal (primitive) operating system can then run on top of the virtual machine
- Created and managed by **Hypervisor**
 - Also known as **Virtual Machine Monitor (VMM)**
 - Two classes of implementations shown next

Type 1 Hypervisor



- Type 1 hypervisor:
 - Provides individual virtual machines to guest OSes
 - eg. IBM VM/370, VMware ESXi

Type 2 Hypervisor



- Type 2 hypervisor OS
 - Runs in host OS
 - Guest OS runs inside Virtual Machine
 - e.g. VMware Workstation, VirtualBox, QEMU

Summary

- Definition of Operating System
- Roles of Operating System
- Common Operating System families
- Operating System structure

Reference

- Modern Operating System (4th Edition)
 - By Andrew S.Tanenbaum
 - Published by Pearson

- Operating System Concepts (8th Edition)
 - By Abraham Silberschatz, Peter Baer Galvin & Greg Gagne
 - Published by McGraw Hill

FYI: MODERN OS FAMILY

(AS OF 2026)

Modern OS: An Overview

- Several dominant desktop OSes:
 - Microsoft Windows family
 - Unix and its variants
 - Mac OS family
- Specialized OSes:
 - Real-time OS
 - Embedded System OS
 - Mobile OS
 - Distributed OS

Microsoft OS Family

- 16-bit:
 - MS-DOS (various versions, v1.0 in 1985)
 - Windows 1.X – 3.X, Windows 9X, Windows ME (2000)
- 32-bit:
 - Windows NT (32-bit, v3.1 in 1994)
 - Windows 2000, XP, 2003, Vista, 7, 8, 10 (2015)
- 64-bit:
 - Windows XP (2005), Vista, 7, 8, 10 (2015), 11 (2021)
- Mostly on PC (Intel / AMD Processors) platforms
- Proprietary
 - some sources available under conditions
- Complex architecture, internals info not widely available


Microsoft OS: Complexity

- Win NT 3.1: (shipped 1993)
 - Dev Team Size: 200
 - 6 Millions LOC (Line of Code)
 - Complete build time is 5 hours on ~486/50
- Win 2000: (shipped 1999)
 - Dev Team Size: 1400
 - 29 Millions LOC (about 50Gb of disk space)
 - Complete build time: 8 hours on 4-way PIII Xeon 550 with 50Gb disk and 512k RAM
- Windows 7: (shipped 2009)
 - Dev Team Size: ~2500 (split into 25 teams of ~100)
 - ~40 Millions LOC

Unix and its Variants

- Many Unix Variants:
 - Unix System V versions
 - Berkeley System Distribution (BSD)
 - Sun Solaris
 - SGI IRIX, IBM AIX, Digital Unix, HP-UX, ...
 - Linux
 - MacOS X (BSD + Mach + Apple)
- Programming Interface (API) mostly the same, fundamentals are same but small differences exist
- Simple architecture, internals well understood, good documentation, research papers
- Linux + BSD open source
- Non-proprietary, POSIX standards
- Implementations on many processors architectures:
 - x86, powerpc, m68k, mips, arm, sparc, alpha, ...


Unix Variants: Rough timeline



Linux Distributions (Distros)

- Most popular form of linux OS
 - E.g. Ubuntu, Debian, Arch Linux, Fedora, BSD, CentOS, etc
- Essentially a software collection with:
 - Linux kernel
 - Common software:
 - Desktop management, browser, media player etc
 - Development tools/libraries
 - gnu libraries, compilers etc
 - Device drivers
 - Package manager

MacOS Timeline



Characteristics of Modern PC OS

- Multitasking:
 - Concurrent execution of programs
 - On multiple cores
- Multi/single user:
 - Unix usually multi-user
 - Windows usually single-user at a time
- Wide range of computer hardware:
 - Single PC/notebook
 - Shared memory systems with 10-100s of processors
 - Machine clusters with 100-1000s of processors
 - Distributed computing on Internet with >10K machines

Real Time OS

- OS for computer systems with time constraints:
 - periodicity, deadlines
- Examples:
 - **Critical Systems**: aircraft flight system, nuclear power plant, radar system
 - **Consumer appliances**: mobile phones, mp3, video players
- **Hard real time**:
 - Timing requirements **must** be respected, eg. control system
- **Soft real time**:
 - Some timing constraints can be missed
- May need **formally verified systems**

Embedded OS

- OS for specialized devices and appliances
- Examples:
 - (Older)Smartphones, microwave oven, car, smart cards, game consoles, etc
- Special consideration required for:
 - power usage, real-time requirement and memory limitations
- Usually:
 - Not general purpose: cannot run any application
 - Cannot be modified: Mostly stored in Read-Only Memory (ROM)

Mobile OS

- OS for smartphones, tablets, PDAs, or other mobile devices
- Examples:
 - Android, iOS, etc
- Characteristics:
 - Essentially a customised version of PC OS
 - Common features: Touchscreen, Cellular, (Video) Camera, etc...

Distributed OS

- OS for computers/processors connected using network
 - Loosely or Tightly coupled
- Loosely coupled:
 - Autonomous nodes, network may be wide area
 - Communication is asynchronous
 - Reliability issues:
 - communication may not be reliable, nodes may fail
 - Resources are distributed
 - eg. distributed filesystem, P2P storage
- Tightly coupled:
 - Specialized node (e.g. Computing nodes) that shares other resources (e.g. Memory / Harddisk etc), nodes in close proximity
 - Examples:
 - Tembusu2 cluster compute nodes:
 - ~100 Intel Xeon nodes running CentOS